

EVERY ROW LEAVES A WAKE.

WAKE Environment Setup and Reproduction Guide

A clean-environment path to verify the public evidence, run the coach dashboard, and reproduce the saved evaluation at US\$0.00.

Prepared for	Hackathon judges, reviewers, and technical evaluators
Submission state	Reproducible replay and owner-approved final media ready
Evidence date	31 August 2026

Evidence-backed rowing review. Replay-first operation. Human-approved memory.

Choose the correct execution path

Mode	Model call	API key	Expected model cost
Clean reproduction	No	Not required	US\$0.00
Dashboard replay	No	Not required	US\$0.00
Optional live	Yes	Required	Authorized per run

Recommended judge path: run the clean reproduction, start replay, inspect the saved dashboard and evaluation, then use live mode only if a new paid investigation is desired.

What this guide reproduces

WAKE - Agentic Rowing Intelligence reconstructs fragmented rowing evidence, routes uncertainty to the right human, and preserves approved rowing memory. The repository ships public fixtures, saved model outputs, verifiers, a replay dashboard, and the exact evaluation artifacts used in the submission.

The default reproduction is intentionally offline with respect to model calls. It verifies committed artifacts and never spends API budget, even if a local key exists.

- 52 pre-regatta and 50 post-regatta synthetic activities, grounded in real rowing formats and plausible value ranges.
- Ten fixed WAKE-versus-baseline evaluation cases and observable saved trajectories.
- A React/Vinext coach interface with sessions, team memory, goal memory, competition review, and evaluation.
- Replay-first execution; live analysis remains an explicit, separately authorized action.

Expected clean verification cost: US\$0.00. Expected live cost: shown before and after each authorized run.

Clean-environment setup

Start from the repository root on macOS or Linux. The locked workflow expects Python 3.11 or newer, uv, and Node.js 22 or newer. Node.js 24.19.0 and npm 10.8.2 were used for the final local audit.

Do not add a .env file for replay. Secrets and private athlete material are intentionally excluded from the source package.

```
unzip wake-source-submission.zip
cd micro1-hackaton
python3 --version
uv --version
node --version
npm --version

./scripts/reproduce_submission.sh
```

If the web build reports that Node is too old, switch to Node 22+ before repeating the command.

Replay without an API key

The full verification path installs locked dependencies, validates public data, checks saved outputs and trajectories, runs Python and web tests, lints the interface, and creates a production build.

When dependencies already exist, use the verification-only path. Both commands ignore OPENAI_API_KEY and make no provider request.

- Complete Git checkout status: READY with repository_ready and final_video_ready both true.
- Extracted source-only ZIP status: PENDING_FINAL_VIDEO with repository_ready true, because the accepted MP4 is uploaded separately.
- Expected fixed-case counts: 10 WAKE outputs, 10 baseline outputs, and 10 trajectories.
- Expected official scores: WAKE 83.76 and direct baseline 49.00.

```
./scripts/reproduce_submission.sh --verify-only  
.  
venv/bin/python scripts/verify_submission_readiness.py --json
```

Run the dashboard in replay mode

Replay is the product default. It opens saved evidence and deterministic club packages without calling a model. The launcher starts the local API and web interface together and stops both with Ctrl+C.

- Sessions: club pulse, attention, crews, athletes, intelligence, and saved reviews.
- Goal memory: coach-approved session memory plus the saved 102-activity club synthesis.
- Competition: connected fictional race outcomes with explicit non-causal boundaries.
- Evaluation: committed WAKE-versus-baseline evidence; opening it costs US\$0.00.

```
./scripts/start_dashboard.sh  
open http://localhost:3000/
```

Optional live investigation

Live mode is not required to judge the saved submission. Use it only when a reviewer wants to execute a new bounded investigation. The key remains local and must never be committed or included in the ZIP.

Every live execution requires a positive finite operational authorization in the interface. That authorization is a start gate, not a provider billing cap. The resulting output displays observed tokens, runtime, and approximate cost.

```
cp .env.example .env  
# Add OPENAI_API_KEY locally; never commit .env  
./scripts/start_dashboard.sh --live
```

A live run may spend money. Replay, reopening saved results, and the fixed reproduction script do not.

Baseline and evaluation

The official comparison holds the model family, ten case summaries, output schema, and grader constant. The direct baseline receives the compact case and produces one structured answer. WAKE receives the same case but can call deterministic tools and must pass a verifier.

The comparison measures the submitted workflows, not superiority over a qualified human coach. The separate club-memory audit is a structural capability check and is not a semantic coaching-quality score.

- WAKE: 83.76/100 across ten cases.
- Direct baseline: 49.00/100 across the same ten cases.
- Incremental official comparison cost: US\$0.283344.
- Preserved regression: environmental interpretation moved from 80.00 to 76.00.

```
./venv/bin/python scripts/run_official_evaluation.py --help  
.  
venv/bin/python scripts/verify_submission_readiness.py
```

Expected outputs and cost

A successful clean run ends with passing public verifiers, Python tests, web tests, lint, and a production build. Saved results can be reopened and regraded without a new model call.

The exact number of tests can grow as the repository evolves; treat a non-zero exit code or failed verifier as a reproduction failure. Review the generated JSON summary for machine-readable evidence counts and scores.

- Replay and clean verification: US\$0.00 model cost.
- Saved 102-activity club memory reopen: US\$0.00.
- Saved club-memory creation: US\$0.037384, 6,322 tokens.
- Three final owner live-QA runs: US\$0.283834 observed total.

Troubleshooting and support boundary

Use the commands below before changing code. They expose the launch plan, verify repository readiness, and show whether the local ports are already occupied.

- Port in use: stop the earlier WAKE launcher or select different --web-port and --api-port values.
- No API key: remain in replay; live mode is intentionally unavailable.
- Private state: keep it under private-data/ or another ignored local path.
- Browser cache: reload after changing the build or restarting the service.

```
./scripts/start_dashboard.sh --print-plan
.venv/bin/python scripts/verify_submission_readiness.py --json
lsof -nP -iTCP:3000 -sTCP:LISTEN
lsof -nP -iTCP:8788 -sTCP:LISTEN
```

This prototype storage is user-restricted but is not encrypted, authenticated, backed up, multi-tenant, or a production database.